



Faster Python Programs - Measure, don't Guess

A Tutorial at PyCon 2018

May 9, 2018

Cleveland, OH, USA

author: Dr.-Ing. Mike Müller
email: mmueller@python-academy.de
twitter: @pyacademy
version: 4.0

Python Academy - The Python Training Specialist

- Dedicated to Python training since 2006
- Wide variety of Python topics
- Experienced Python trainers
- Trainers are specialists and often core developers of taught software
- All levels of participants from novice to experienced software developers
- Courses for system administrators
- Courses for scientists and engineers
- Python for data analysis
- Open courses around Europe
- Customized in-house courses
- Python programming
- Development of technical and scientific software
- Code reviews
- Coaching of individuals and teams migrating to Python
- Workshops on developed software with detailed explanations

More Information: www.python-academy.com

Current Training Modules - Python Academy

As of 2018

Module Topic	Length (days)	In-house	Open
Python for Programmers	3	yes	yes
Python for Non-Programmers	4	yes	yes
Advanced Python	3	yes	yes
Python for Scientists and Engineers	3	yes	yes
Python for Data Analysis	3	yes	yes
Machine Learning with Python	3	yes	no
Professional Testing with Python	3	yes	yes
High Performance Computing with Python	3	yes	yes
Cython in Depth	2	yes	yes
Introduction to Django	4	yes	yes
Advanced Django	3	yes	yes
SQLAlchemy	1	yes	yes
High Performance XML with Python	1	yes	yes
Optimizing Python Programs	1	yes	yes
Python Extensions with Other Languages	1	yes	no
Data Storage with Python	1	yes	no
Introduction to Software Engineering with Python	1	yes	no
Introduction to wxPython	1	yes	no
Introduction to PySide/PyQt	1	yes	no
Overview of the Python Standard Library	1	yes	no
Threads and Processes in Python	1	yes	no
Windows Programming with Python	1	yes	no
Network Programming with Python	1	yes	no
Introduction to IronPython	1	yes	no

We offer on-site and open courses world-wide. We customize and extend training modules for your needs. We also provide consulting services such as code review, customized programming, and workshops.

More information: www.python-academy.com

Contents

1	How Fast is Fast Enough?	6
1.1	Introduction	6
1.2	Optimization Guidelines	6
2	Strategy	8
2.1	Measuring in Stones	8
2.2	Timing	9
2.3	Profiling CPU Usage	10
2.3.1	Working with Standard Python	11
2.3.2	Working with IPython and Jupyter Notebook	14
2.4	Wall Clock vs. CPU Time	16
2.5	A More Complex Function	21
2.6	A Picture is Worth a Thousand Words	23
2.6.1	SnakeVis	23
2.7	Going Line-by-Line	24
2.7.1	Exercise	32
2.8	Profiling Memory Usage	33
2.8.1	Pympler	33
2.8.2	Memory Usage Line-by-Line with <code>memory_profiler</code>	40
2.8.3	Exercise	43
3	Algorithms and Anti-patterns	44
3.1	String Concatenation	44
3.1.1	Exercise	44
3.2	List and Generator Expressions	44
3.2.1	Exercise	45
3.3	Think Global buy Local	45
4	The Right Data Structure	48
4.1	Use built-in Data Types	48
4.2	<code>list</code> vs. <code>set</code>	48
4.3	<code>list</code> vs. <code>deque</code>	59
4.4	<code>dict</code> vs. <code>defaultdict</code>	63
4.5	Big-O notation and Data Structures	66
4.6	$O(1)$ vs. $O(n)$ vs. $O(n^2)$	67

4.7 Exercises	69
5 Caching	70
5.1 Reuse before You Recalculate	70
5.2 Deterministic caching	70
5.3 Non-deterministic caching	72
5.4 Least Recently Used Cache	74
5.5 Memcached	75
6 Compilation of Tools for Speedup and Extending	76
7 End	77
7.1 Colophon	77
7.2 Links	77

1 How Fast is Fast Enough?

1.1 Introduction

Since Python is an interpreted language, some types of computations are slower in Python than in compiled languages. Depending on the application, this may or may not be a problem. This tutorial introduces several methods to speed up Python. Before starting to optimize, however the cost involved should be considered. Optimized code may need more effort to develop and maintain, leading to prolonged development time. So there is always a balance between speed of development and speed of program execution.

1.2 Optimization Guidelines

Premature optimization is the root of all evil.

D. Knuth or C. A. R. Hoare or folklore (attributions seems not totally clear)

Before you start thinking about optimization make sure your program works correctly. Never optimize before the program produces the desired results.

Optimization often comes with a price: It tends to make your code less readable. Since most of the programming time for software is spent on maintenance rather than developing new code, readability and maintainability is of great importance for an effective life cycle of your program. Therefore, always think twice if it is really worth before you make your code less readable the speed gain. After all, we deliberately choose Python for its excellent readability and pay with somewhat slower programs for certain tasks.

A few general guidelines are formulated as follows:

1. Make sure your program is really too slow. Do you really need more performance? Are there any other slowdown factors such as network traffic or user input that have more impact on speed? Does it hurt if the program runs slowly?
2. Don't optimize as you go. Don't waste time before you are certain that you will need the additional speed.
3. Only realistic use cases and user experience should be considered.
4. Architecture can be essential for performance. Is it appropriate?
5. Are there any bugs that slow down the program?
6. If the program is really too slow, find the bottlenecks by profiling (use module `profile`).
7. Always check the result of optimization with all unit tests. Don't optimize with bugs.

Usually for complex programs, the most performance gain can be achieved by optimization of algorithms. Finding what big-O notation an algorithm has is very important to predict performance for large amounts of data.

The first thing you should check before making any changes in your program is external causes that may slow down your program. Likely candidates are:

- network connections
- database access
- calls to system functions

1 How Fast is Fast Enough?

In most cases hardware is cheaper than programmer time. Always check if there is enough memory for application. Swapping memory pages to disc may slow down execution by an order of magnitude. Make sure you have plenty of free disk space and a recent and fast processor. The Python Cookbook [MART2005], also available [online](#)¹, is a very good compilation of short and not so short solutions to specific problems. Some of the recipes, especially in the algorithm section are applicable to performance issues.

The Python in a Nutshell book ([MART2006]) contains a good summary on optimization, including profiling as well as large-scale and small-scale optimization (see pages 474 to 489). There are two chapters about optimization in [ZIAD2008]. A good resource for scientific applications in Python is [LANG2006] that also contains substantial material on optimization and extending of Python with other languages.

Some of them are exemplified in the following section.

From now on we assume you have done all the above-mentioned steps and still need more speed.

2 Strategy

2.1 Measuring in Stones

Programs will run at different speeds on different hardware. The use of benchmarks allows to measure how fast your hardware and, in the case of Python, how fast the used implementation is. Python has the module `test.pystone` that allows to benchmark hardware and implementation. We can use it as a standalone script.

With Python 2.7:

```
python pystone2_3.py
Pystone(1.1) time for 50000 passes = 0.303863
This machine benchmarks at 164548 pystones/second

python pystone2.py
Pystone(1.1) time for 50000 passes = 0.311293
This machine benchmarks at 160620 pystones/second
```

Python 3.5 is a little bit slower:

::

```
(py35) python pystone2_3.py Pystone(1.1) time for 50000 passes = 0.365526 This machine benchmarks
at 136789 pystones/second
```

Python 3.6 is about as fast as Python 3.5:

```
(py36) python pystone2_3.py
Pystone(1.1) time for 50000 passes = 0.363429
This machine benchmarks at 137578 pystones/second
```

PyPy is **significantly** faster than CPython for this benchmark:

```
pypy pystone2_3.py
Pystone(1.1) time for 50000 passes = 0.034823
This machine benchmarks at 1.43583e+06 pystones/second
```

But Jython is slower than CPython:

```
jython2.7 pystone2_3.py
Pystone(1.1) time for 50000 passes = 0.47303
This machine benchmarks at 105701 pystones/second
```

We can also use `pystone` in our programs:

```
>>> from test import pystone
>>> pystone.pystones()
(1.2585885652668103, 39727.04136987008)
```


2.2 Timing

The first value is the benchmark time in seconds and the second the pystones. We can use the pystone value to convert measured run times in seconds into pystones:

```
# file: pystone_converter.py

"""Convert seconds to kilo pystones."""

from test import pystone

BENCHMARK_TIME, PYSTONES = pystone.pystones()

def kpystone_from_seconds(seconds):
    """Convert seconds to kilo pystones."""
    return (seconds * PYSTONES) / 1e3

if __name__ == '__main__':

    def test():
        """Show how it works"""

        print
        print '%10s %10s' % ('seconds', 'kpystones')
        print
        for seconds in [0.1, 0.5, 1.0, 2.0, 5.0]:
            print ('%10.5f %10.5f' % (seconds, kpystone_from_seconds(seconds)))

    test()
```

We will use this function to compare our results.

2.2 Timing

Python provides the module `timeit` to measure how long the execution of code takes. It repeats it a specific times to get a reasonable measurements. Let's have a look at a small example:

```
import timeit

a = 1
b = 2

number = 10**6 # defaults to 100,000

total_duration = timeit.timeit('a + b', globals=locals(), number=number)
print(total_duration / number)
```

prints:

```
5.795985506847501e-08
```

2.3 Profiling CPU Usage

The Jupyter Notebook has very useful magic command `%timeit` that makes the use much simple. We can achieve the same (and a bit more) with:

```
%timeit a + b
```

which prints:

51.6 ns \pm 0.338 ns per loop (mean \pm std. dev. of 7 runs, 10000000 loops each)

`%timeit` repeats determines the number of loops so that it takes about 1 second for all loops. The loops are multiples of 10. Finally, it repeats this process 7 times and shows simple statistics of the results. All the parameters can be adjusted. Use `timeit.timeit?` to see all options.

As always `%timeit` with one `%` works on one line and `%%timeit` with two `%` works on a whole cell.

2.3 Profiling CPU Usage

There are two modules in the Python standard library that allow measuring the used CPU time:

- `profile`
- `cProfile`

Because `profile` is a pure Python implementation that adds a lot of overhead, `cProfile` is the recommended tool. Up to Python 2.7 there was another pure Python profiler `hotshot`. It has been removed from Python 3. These profilers are deterministic and therefore actually run the code they are profiling and measure its execution time. This has some overhead but provides reliable results in most cases. `cProfile` tries to minimize this overhead. The other type of profiling is called statistical and uses random sampling of the effective instruction pointer. This has less overhead but is also less precise. We won't look at statistical techniques here.

Let's write a small program whose whole purpose is to use up CPU time:

```
# file profile_me.py

"""Example to be profiled.
"""
import sys
import time

if sys.version_info.major < 3:
    range = xrange

def fast():
    """Wait 0.001 seconds.
    """
    time.sleep(1e-3)

def slow():
    """Wait 0.1 seconds.
```

2.3 Profiling CPU Usage

```
"""
time.sleep(0.1)

def use_fast():
    """Call `fast` 100 times.
    """
    for _ in range(100):
        fast()

def use_slow():
    """Call `slow` 100 times.
    """
    for _ in range(100):
        slow()

if __name__ == '__main__':
    use_fast()
    use_slow()
```

2.3.1 Working with Standard Python

Now we import our module as well as cProfile:

```
>>> import profile_me
>>> import cProfile
```

and make an instance of Profile:

```
>>> profiler = cProfile.Profile()
```

First we call our fast function:

```
>>> profiler.runcall(profile_me.use_fast)
```

and look at the statistics cProfile provides:

```
>>> profiler.print_stats()
202 function calls in 0.195 CPU seconds
Ordered by: standard name
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
   100    0.000    0.000    0.195    0.002 profile_me.py:3(fast)
     1    0.000    0.000    0.195    0.195 profile_me.py:9(use_fast)
     1    0.000    0.000    0.000    0.000 ~:0(<method 'disable' of
```

2.3 Profiling CPU Usage

```
'_lsprof.Profiler' objects>)\n100    0.194    0.002    0.194    0.002 ~:0(<time.sleep>)
```

The column headers have the following meaning:

- `ncalls` is the number of calls to this function
- `totttime` is the total time spent in this function, where calls to sub-functions are excluded from time measurement
- `percall` is `totttime` divided by `ncalls`
- `cumtime` is the cumulative time, that is the total time spent in this including the time spent in sub-functions
- `percall` is `cumtime` divided by `ncalls`
- `filename:lineno(function)` are the name of the module, the line number and the name of the function

We can see that the function `fast` is called 100 times and that it takes about 0.002 seconds per call. At first look it is surprising that `totttime` is zero. But if we look at the time the function `time.sleep` uses up, it becomes clear the `fast` spends only 0.001 seconds (0.195 - 0.194 seconds) and the rest of the time is burnt in `time.sleep()`.

We can do the same thing for our slow function:

```
>>> profiler = cProfile.Profile()\n>>> profiler.runcall(profile_me.use_slow)\n>>> profiler.print_stats()\n      202 function calls in 10.058 CPU seconds\nOrdered by: standard name\nncalls  tottime  percall  cumtime  percall filename:lineno(function)\n      1   0.001   0.001   10.058   10.058 profile_me.py:13(use_slow)\n     100   0.001   0.000   10.058   0.101 profile_me.py:6(slow)\n      1   0.000   0.000   0.000   0.000 ~:0(<method 'disable' of\n'_lsprof.Profiler' objects>)\n     100  10.057   0.101   10.057   0.101 ~:0(<time.sleep>)
```

Not surprisingly, the run times are nearly two orders of magnitude greater, because we let `sleep` use up one hundred times more time.

Another method to invoke the profiler is to use the function `run`:

```
>>> cProfile.run('profile_me.use_fast()')\n      203 function calls in 0.195 CPU seconds\nOrdered by: standard name\nncalls  tottime  percall  cumtime  percall filename:lineno(function)\n      1   0.000   0.000   0.195   0.195 <string>:1(<module>)\n     100   0.000   0.000   0.195   0.002 profile_me.py:3(fast)\n      1   0.000   0.000   0.195   0.195 profile_me.py:9(use_fast)\n      1   0.000   0.000   0.000   0.000 ~:0(<method 'disable' of
```

2.3 Profiling CPU Usage

```
'_lsprof.Profiler' objects>)\n100    0.195    0.002    0.195    0.002 ~:0(<time.sleep>)
```

Here we supply the function to be called as a string with parenthesis, i.e. a string that can be used in an `exec` statement as opposed to the function object we supplied to the `runcall` method of our `Profile` instance.

We can also supply a file where the measured runtime data will be stored:

```
>>> cProfile.run('profile_me.use_fast()', 'fast.stats')
```

Now we can use the `pstats` module to analyze these data:

```
>>> cProfile.run('profile_me.use_fast()', 'fast.stats')\n>>> import pstats\n>>> stats = pstats.Stats('fast.stats')
```

We can just print out the data in the same format we saw before:

```
>>> stats.print_stats()\nWed Mar 11 16:11:39 2009      fast.stats\n      203 function calls in 0.195 CPU seconds\n      Random listing order was used\nncalls  tottime  percall  cumtime  percall filename:lineno(function)\n   100     0.194     0.002     0.194     0.002 ~:0(<time.sleep>)\n     1     0.000     0.000     0.000     0.000 ~:0(<method 'disable' of '_lsprof.Profiler'\nobjects>)\n   100     0.000     0.000     0.195     0.002 profile_me.py:3(fast)\n     1     0.000     0.000     0.195     0.195 <string>:1(<module>)\n     1     0.000     0.000     0.195     0.195 profile_me.py:9(use_fast)
```

We can also sort by different columns and restrict the number of lines printed out. Here we sort by the number of calls and want to see only the first three columns:

```
>>> stats.sort_stats('calls').print_stats(3)\nWed Mar 11 16:11:39 2009      fast.stats\n      203 function calls in 0.195 CPU seconds\n      Ordered by: call count\n      List reduced from 5 to 3 due to restriction <3>\nncalls  tottime  percall  cumtime  percall filename:lineno(function)\n   100     0.194     0.002     0.194     0.002 ~:0(<time.sleep>)\n   100     0.000     0.000     0.195     0.002 profile_me.py:3(fast)\n     1     0.000     0.000     0.000     0.000 ~:0(<method 'disable' of\n'_lsprof.Profiler' objects>)
```

Or we sort by time used and show all lines:

2.3 Profiling CPU Usage

```
>>> stats.sort_stats('time').print_stats()
Wed Mar 11 16:11:39 2009      fast.stats
      203 function calls in 0.195 CPU seconds
Ordered by: internal time
ncalls  tottime  percall  cumtime  percall filename:lineno(function)
   100    0.194    0.002    0.194    0.002 ~:0(<time.sleep>)
   100    0.000    0.000    0.195    0.002 profile_me.py:3(fast)
     1    0.000    0.000    0.195    0.195 profile_me.py:9(use_fast)
     1    0.000    0.000    0.195    0.195 <string>:1(<module>)
     1    0.000    0.000    0.000    0.000 ~:0(<method 'disable' of
      '_lsprof.Profiler' objects>)
```

We can also get information about which function is called by a certain function:

```
>>> stats.print_callers('fast')
Ordered by: internal time
List reduced from 5 to 2 due to restriction <'fast'>
Function                                was called by...
profile_me.py:3(fast)                    profile_me.py:9(use_fast)
((100, 100, 0.00040628818660897912, 0 .19478914258667296))    0.195
profile_me.py:9(use_fast)                <string>:1(<module>)
((1, 1, 0.00026121123840443721, 0.1950503538250774))    0.195
```

We can also find out what functions are called:

```
>>> stats.print_callees('use_fast')
Ordered by: internal time
List reduced from 5 to 1 due to restriction <'use_fast'>
Function                                called...
profile_me.py:9(use_fast)                profile_me.py:3(fast)
((100, 100, 0.00040628818660897912, 0.19478914258667296))    0.195
```

There are more interesting attributes such as the number of calls:

```
>>> stats.total_calls
203
```

2.3.2 Working with IPython and Jupyter Notebook

In IPython you can get similar results as with `profile.run()` with the magic `%prun`:

```
%prun profile_me.use_fast()
```

```
204 function calls in 0.114 seconds

Ordered by: internal time
```

2.3 Profiling CPU Usage

```
ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
   100    0.113    0.001    0.113    0.001  {built-in method sleep}
   100    0.001    0.000    0.114    0.001  profile_me.py:12(fast)
     1    0.000    0.000    0.114    0.114  profile_me.py:24(use_fast)
     1    0.000    0.000    0.114    0.114  {built-in method exec}
     1    0.000    0.000    0.114    0.114  <string>:1(<module>)
     1    0.000    0.000    0.000    0.000  {method 'disable' of '_lsprof.Profiler' objects}
```

You can limit the number of output lines:

```
%prun -l 2 profile_me.use_fast()

204 function calls in 0.114 seconds
```

```
Ordered by: internal time
List reduced from 6 to 2 due to restriction <2>

ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
   100    0.113    0.001    0.113    0.001  {built-in method sleep}
   100    0.001    0.000    0.113    0.001  profile_me.py:12(fast)
```

Using a string as a filter, our output shows only functions with this string in their names. In our case word fast:

```
%prun -l fast profile_me.use_fast()

204 function calls in 0.113 seconds

Ordered by: internal time
List reduced from 6 to 2 due to restriction <'fast'>

ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
   100    0.001    0.000    0.114    0.001  profile_me.py:12(fast)
     1    0.001    0.001    0.115    0.115  profile_me.py:24(use_fast)
```

We can create an `pstats.Stats` object:

```
stats = %prun -r profile_me.use_fast()
```

That behaves as the one generated with standard Python:

```
stats.sort_stats('calls').print_stats(3)
```

```
204 function calls in 0.113 seconds
```

2.4 Wall Clock vs. CPU Time

```
Ordered by: call count
List reduced from 6 to 3 due to restriction <3>

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
   100    0.001    0.000    0.113    0.001 profile_me.py:12(fast)
   100    0.112    0.001    0.112    0.001 {built-in method sleep}
     1    0.000    0.000    0.113    0.113 {built-in method exec}
```

The option `-s` sorts and works multiple times:

```
%prun -s calls -s time profile_me.use_fast()
```

```
204 function calls in 0.113 seconds

Ordered by: call count, internal time

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
   100    0.113    0.001    0.113    0.001 {built-in method sleep}
   100    0.000    0.000    0.113    0.001 profile_me.py:12(fast)
     1    0.000    0.000    0.113    0.113 profile_me.py:24(use_fast)
     1    0.000    0.000    0.113    0.113 {built-in method exec}
     1    0.000    0.000    0.113    0.113 <string>:1(<module>)
     1    0.000    0.000    0.000    0.000 {method 'disable' of '_lsprof.Profiler' objects}
```

You can also save your results to a file. Either the text as seen on the screen:

```
%prun -T stats.txt profile_me.use_fast()
```

```
*** Profile printout saved to text file 'stats.txt'.
```

now `stats.txt` contains the screen output.

Or, you dump it into a stats file for later use with other tools:

```
%prun -D ipython.stats profile_me.use_fast()
```

```
*** Profile stats marshalled to file 'ipython.stats'.
```

2.4 Wall Clock vs. CPU Time

Per default `cProfile` measures wall clock time, i.e. the time elapsed between start and end of the function. Since typically computers do more than one thing at a time, this times usually does not correspond with the usage time of the CPU. Also, often computers often wait for IO. During this time the CPU is more or less idle but the time elapses nevertheless.

2.4 Wall Clock vs. CPU Time

Unfortunately, there are differences between operating systems in how they measure CPU time. Let's look at a simple test function:

```
# file: measuring/clock_check.py

"""Checking different timing functions.
"""

from __future__ import print_function

import os
import sys
import time
import timeit

if sys.version_info.major < 3:
    range = xrange

def clock_check(duration=1):
    """Check the measured time with different methods.
    """
    start_os_time0 = os.times()[0]
    start_time_clock = time.clock()
    start_default_timer = timeit.default_timer()
    for _ in range(int(1e6)):
        1 + 1
    time.sleep(duration)
    durtation_os_time0 = os.times()[0] - start_os_time0
    durtation_time_clock = time.clock() - start_time_clock
    durtation_default_timer = timeit.default_timer() - start_default_timer
    print('durtation_os_time0:      ', durtation_os_time0)
    print('durtation_time_clock:    ', durtation_time_clock)
    print('durtation_default_timer:', durtation_default_timer)

if __name__ == '__main__':
    clock_check()
```

We use three different methods to get time stamps:

1. `os.times()[0]` provides the CPU time on all operating systems. While it has six decimals, i.e. microseconds accuracy on Windows, it is only two significant decimals on Unix-like systems.
2. `start_time_clock = time.clock()` is the CPU time on Unix but the wall clock time on Windows.
3. `timeit.default_timer()` chooses the right timing function for wall clock, i.e. `time.time()` on Unix and `time.clock()` on windows.

This is our output on Unix/Linux/MacOSX:

2.4 Wall Clock vs. CPU Time

```
durtation_os_time0:      0.05
durtation_time_clock:    0.041949
durtation_default_timer: 1.04296183586
```

and on Windows:

```
durtation_os_time0:      0.03125
durtation_time_clock:    1.02673477293
durtation_default_timer: 1.02673567261
```

We write a script to look at how `cProfile` can be used to measure both, wall clock and CPU time:

```
# file: cpu_time.py

"""Measuring CPU time instead of wall clock time.
"""

import cProfile
import os
import sys
import time

# Make it work with Python 2 and Python 3.
if sys.version_info.major < 3:
    range = xrange
```

After some imports and a Python 2/3 compatibility helper, we define a function to measure CPU time that is aware of the differences between operating systems:

```
def cpu_time():
    """Function for cpu time. Os dependent.
    """
    if sys.platform == 'win32':
        return os.times()[0]
    else:
        return time.clock()
```

We use two functions:

```
def sleep():
    """Wait 2 seconds.
    """
    time.sleep(2)

def count():
    """100 million loops.
```

2.4 Wall Clock vs. CPU Time

```
"""
for _ in range(int(1e8)):
    1 + 1

def test():
    """Run functions
    """
    sleep()
    count()
```

One that sleeps and one that just loops many times to consume CPU time. We put them into a test function and use `cProfile` with different timing methods:

```
def profile():
    """Profile with wall clock and cpu time.
    """
    profiler = cProfile.Profile()
    profiler.run('test()')
    profiler.print_stats()

    profiler = cProfile.Profile(cpu_time)
    profiler.run('test()')
    profiler.print_stats()

if __name__ == '__main__':
    profile()
```

Without providing a time measurement function: `profiler = cProfile.Profile()`, we get the default wall clock timing.

Providing our timer function: `cProfile.Profile(cpu_time)`, we get the CPU time.

The output on Windows:

```
$ cpu_time.py
    6 function calls in 5.233 seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
      1     0.010      0.010      5.233      5.233  <string>:1(<module>)
      1     0.000      0.000      2.000      2.000  cpu_time.py:25(sleep)
      1     3.222      3.222      3.222      3.222  cpu_time.py:31(count)
      1     0.000      0.000      5.222      5.222  cpu_time.py:38(test)
      1     0.000      0.000      0.000      0.000  {method 'disable' of '_lsprof.Profiler' object}
      1     2.000      2.000      2.000      2.000  {time.sleep}
```

2.4 Wall Clock vs. CPU Time

```
6 function calls in 3.141 seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
1      0.000    0.000    3.141    3.141 <string>:1(<module>)
1      0.000    0.000    0.000    0.000 cpu_time.py:25(sleep)
1      3.141    3.141    3.141    3.141 cpu_time.py:31(count)
1      0.000    0.000    3.141    3.141 cpu_time.py:38(test)
1      0.000    0.000    0.000    0.000 {method 'disable' of '_lsprof.Profiler' object}
1      0.000    0.000    0.000    0.000 {time.sleep}
```

And the output on Unix/Linux/MacOSX:

```
$ python cpu_time.py
6 function calls in 7.171 seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
1      0.000    0.000    7.171    7.171 <string>:1(<module>)
1      0.000    0.000    2.001    2.001 cpu_time.py:25(sleep)
1      5.169    5.169    5.169    5.169 cpu_time.py:31(count)
1      0.000    0.000    7.171    7.171 cpu_time.py:38(test)
1      0.000    0.000    0.000    0.000 {method 'disable' of '_lsprof.Profiler' object}
1      2.001    2.001    2.001    2.001 {time.sleep}

6 function calls in 4.360 seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
1      0.000    0.000    4.360    4.360 <string>:1(<module>)
1      0.000    0.000    0.000    0.000 cpu_time.py:25(sleep)
1      4.360    4.360    4.360    4.360 cpu_time.py:31(count)
1      0.000    0.000    4.360    4.360 cpu_time.py:38(test)
1      0.000    0.000    0.000    0.000 {method 'disable' of '_lsprof.Profiler' object}
1      0.000    0.000    0.000    0.000 {time.sleep}
```

Both seem to correspond.

We can also time with Python 3:

```
$ python3 cpu_time.py
7 function calls in 6.223 seconds

Ordered by: standard name
```

2.5 A More Complex Function

```
ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
      1     0.000     0.000     6.223     6.223  <string>:1(<module>)
      1     0.000     0.000     2.000     2.000  cpu_time.py:25(sleep)
      1     4.223     4.223     4.223     4.223  cpu_time.py:31(count)
      1     0.000     0.000     6.223     6.223  cpu_time.py:38(test)
      1     0.000     0.000     6.223     6.223  {built-in method exec}
      1     2.000     2.000     2.000     2.000  {built-in method sleep}
      1     0.000     0.000     0.000     0.000  {method 'disable' of '_lsprof.Profiler' object}
```

7 function calls in 4.231 seconds

Ordered by: standard name

```
ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
      1     0.000     0.000     4.231     4.231  <string>:1(<module>)
      1     0.000     0.000     0.000     0.000  cpu_time.py:25(sleep)
      1     4.230     4.230     4.230     4.230  cpu_time.py:31(count)
      1     0.000     0.000     4.231     4.231  cpu_time.py:38(test)
      1     0.000     0.000     4.231     4.231  {built-in method exec}
      1     0.000     0.000     0.000     0.000  {built-in method sleep}
      1     0.000     0.000     0.000     0.000  {method 'disable' of '_lsprof.Profiler' object}
```

Conclusion: Always be aware what you are actually measuring. Don't assume to be on particular operating system, try to make your program run cross platform.

2.5 A More Complex Function

We want to explore a bit more fancy function. This is one to calculate π with the Monte Carlo method:

```
# file: simple_pi.py

"""Calculating pi with Monte Carlo.
"""

from __future__ import print_function

import math
import random
import sys

if sys.version_info[0] < 3:
    range = xrange

def pi_plain(total):
    """Calculate pi with `total` hits.
    """
```

2.5 A More Complex Function

```
count_inside = 0
for _ in range(total):
    x = random.random()
    y = random.random()
    dist = math.sqrt(x * x + y * y)
    if dist < 1:
        count_inside += 1
return 4.0 * count_inside / total

if __name__ == '__main__':

    def test():
        """Check if it works.
        """
        n = int(1e6)
        print('pi:', pi_plain(n))

    test()
```

We can also use NumPy for this:

```
# file: numpy_pi.py
"""Calculating pi with Monte Carlo Method and NumPy.
"""

from __future__ import print_function

import numpy #1

def pi_numpy(total): #2
    """Compute pi.
    """
    x = numpy.random.rand(total) #3
    y = numpy.random.rand(total) #4
    dist = numpy.sqrt(x * x + y * y) #5
    count_inside = len(dist[dist < 1]) #6
    return 4.0 * count_inside / total

if __name__ == '__main__':

    def test():
        """Time the execution.
        """
        import timeit
        start = timeit.default_timer()
        pi_numpy(int(1e6))
        print('run time', timeit.default_timer() - start)

    test()
```

2.6 A Picture is Worth a Thousand Words

Doing the statistics with tables is worthwhile and interesting. But there is another way to look at the profiling results: making graphs.

2.6.1 SnakeVis

A nice visualization tool is [SnakeVis](#)². The installation via `pip` is standard. There is a command line version `snakeviz`. Calling it with a name of a file that contains profiling information from the command line will open a browser with the visualization.


We create our profiling statistics at the command line:

```
python -m cProfile -o pi.stat simple_pi.py
```

Now, we start SnakeViz:

```
snakeviz pi.stat
```

This graph will show up in the browser:



The logic of the results is somewhat reversed to that of `RunSnakeRun`. The functions are color-coded. One color represents one function. The main function is the closed circle in the middle. The functions it calls are represented by the arcs outside of it. Hovering over an arc with the mouse, will highlight it and show the portion of the time that is used up in the function itself. The rest will be used by functions it calls. In our example the function `pi()` is the second outer arc. It also takes up the majority of the outermost arc because a lot of the run time is spent in this function itself. The second biggest arc is `random()` followed by `sqrt()`.

Clicking on an arc will reset the display in such a way that the selected function becomes the full circle in the middle and the other, more outer, functions are rearranged accordingly. This works like a zoom.

We can use the NumPy version. Creating the profiling statistics:


2.7 Going Line-by-Line

```
python -m cProfile -o numpy_pi.stats numpy_pi.py
```

Showing the visualization:

```
snakeviz numpy_pi.stats
```

yields a much more complex graph:



Many functions are called by more than one function. This results in a color appearing at many isolated places.

There is also an extension for Jupyter Notebook (IPython). We can install it with

```
%load_ext snakeviz
```

and visualize one line of code:

```
%snakeviz func()
```

or a whole cell with multiple lines of code:

```
%%snakeviz
# code for profiling and visualization
# here
```

2.7 Going Line-by-Line

With `cProfile` the finest resolution we get is the function call. But there is `line_profiler` by Robert Kern that allows line-by-line profiling. `line_profiler` comes bundled with `kernprof` that adds some features to `cProfile`. The installation is simple:

2.7 Going Line-by-Line

```
pip install line_profiler
```

We can use `kernprof` from the command line, which just uses `cProfile`. The option `-v` shows the statistics right away:

```
$ kernprof -v profile_me.py
Wrote profile results to profile_me.py.prof
  406 function calls in 10.204 seconds
```

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	10.204	10.204	<string>:1(<module>)
100	0.001	0.000	10.081	0.101	profile_me.py:15(slow)
1	0.001	0.001	0.121	0.121	profile_me.py:21(use_fast)
1	0.001	0.001	10.082	10.082	profile_me.py:28(use_slow)
1	0.001	0.001	10.204	10.204	profile_me.py:4(<module>)
100	0.001	0.000	0.120	0.001	profile_me.py:9(fast)
1	0.000	0.000	10.204	10.204	{execfile}
1	0.000	0.000	0.000	0.000	{method 'disable' of '_lsprof.Profiler' objects}
200	10.199	0.051	10.199	0.051	{time.sleep}

We add the decorator `profile` to the function we would like to profile:

```
# file profile_me_use_line_profiler.py

"""Example to be profiled.
"""

import time
import sys

if sys.version_info.major < 3:
    range = xrange

def fast():
    """Wait 0.001 seconds.
    """
    time.sleep(1e-3)

def slow():
    """Wait 0.1 seconds.
    """
    time.sleep(0.1)
```

2.7 Going Line-by-Line

```
@profile
def use_fast():
    """Call `fast` 100 times.
    """
    for _ in range(100):
        fast()

@profile
def use_slow():
    """Call `slow` 100 times.
    """
    for _ in range(100):
        slow()

if __name__ == '__main__':
    use_fast()
    use_slow()
```

Now we can use the option `-l` to turn on `line_profiler`:

```
$ kernprof -l -v profile_me_use_line_profiler.py
Wrote profile results to profile_me_use_line_profiler.py.lprof
Timer unit: 1e-06 s

File: profile_me_use_line_profiler.py
Function: use_fast at line 20
Total time: 0.120634 s

Line #      Hits          Time  Per Hit   % Time  Line Contents
=====
    20                               @profile
    21                               def use_fast():
    22                                   """Call `fast` 100 times.
    23                                   """
    24          101           732        7.2     0.6      for _ in range(100):
    25          100      119902     1199.0    99.4          fast()

File: profile_me_use_line_profiler.py
Function: use_slow at line 27
Total time: 10.086 s

Line #      Hits          Time  Per Hit   % Time  Line Contents
=====
    27                               @profile
    28                               def use_slow():
    29                                   """Call `slow` 100 times.
    30                                   """
```

2.7 Going Line-by-Line

31	101	1147	11.4	0.0	for _ in range(100):
32	100	10084845	100848.4	100.0	slow()

This shows us how much time each line used. Our test functions are very short. Let's create a small function that accumulates the sums of all elements in a list:

```
# file accumulate.py

"""Simple test function for line_profiler.
"""

@profile
def accumulate(iterable):
    """Accumulate the intermediate steps in summing all elements.

    The result is a list with the length of `iterable`.
    The last element is the sum of all elements of `iterable`
    >>> accumulate(range(5))
    [0, 1, 3, 6, 10]
    >>> accumulate(range(10))
    [0, 1, 3, 6, 10, 15, 21, 28, 36, 45]
    """
    acm = [iterable[0]]
    for elem in iterable[1:]:
        old_value = acm[-1]
        new_value = old_value + elem
        acm.append(new_value)
    return acm

if __name__ == '__main__':
    accumulate(range(10))
    accumulate(range(100))
```

Let's look at the output:

```
$ kernprof -l -v accumulate.py
Wrote profile results to accumulate.py.lprof
Timer unit: 1e-06 s

File: accumulate.py
Function: accumulate at line 3
Total time: 0.000425 s

Line #      Hits          Time  Per Hit   % Time  Line Contents
=====
      3                 3             @profile
      4                4             def accumulate(iterable):
```

2.7 Going Line-by-Line

```
5                                     """Accumulate the intermediate steps i
6
7                                     The result is a list with the lenght o
8                                     The last elements is the sum of all ele
9                                     >>>accumulate(range(5))
10                                    [0, 1, 3, 6, 10]
11                                    accumulate(range(10))
12                                    [0, 1, 3, 6, 10, 15, 21, 28, 36, 45]
13                                    """
14          2          5          2.5          1.2          acm = [iterable[0]]
15          110         99          0.9          23.3         for elem in iterable[1:]:
16          108         94          0.9          22.1             old_value = acm[-1]
17          108         98          0.9          23.1             new_value = old_value + elem
18          108         127         1.2          29.9             acm.append(new_value)
19          2          2          1.0          0.5             return acm
```

The algorithm could be written more concisely. In fact, the three lines inside the loop could be one. But we would like to see how much each operation takes and therefore spread things over several lines.

Another example looks at some simple mathematical calculations:

```
#calc.py

"""Simple test function for line_profiler doing some math.
"""

import math
import sys

if sys.version_info.major < 3:
    range = xrange

@profile
def calc(number, loops=1000):
    """Do some math calculations.
    """
    sqrt = math.sqrt
    for x in range(loops):
        x = number + 10
        x = number * 10
        x = number ** 10
        x = pow(number, 10)
        x = math.sqrt(number)
        x = sqrt(number)
        math.sqrt
        sqrt
```

2.7 Going Line-by-Line

```
if __name__ == '__main__':  
    calc(100, int(1e5))
```

The output shows which operation takes the most time:

```
$ kernprof -l -v calc.py  
Wrote profile results to calc.py.lprof  
Timer unit: 1e-06 s  
  
File: calc.py  
Function: calc at line 7  
Total time: 1.33158 s  
  
Line #      Hits          Time  Per Hit   % Time  Line Contents  
=====
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
7					@profile
8					def calc(number, loops=1000):
9					"""Do some math calculations.
10					"""
11	1	4	4.0	0.0	sqrt = math.sqrt
12	100001	77315	0.8	5.8	for x in range(loops):
13	100000	87124	0.9	6.5	x = number + 10
14	100000	84518	0.8	6.3	x = number * 10
15	100000	330587	3.3	24.8	x = number ** 10
16	100000	378584	3.8	28.4	x = pow(x, 10)
17	100000	109849	1.1	8.2	x = math.sqrt(number)
18	100000	93211	0.9	7.0	x = sqrt(number)
19	100000	88768	0.9	6.7	math.sqrt
20	100000	81624	0.8	6.1	sqrt

The function `pow` takes by far the most time, whereas `sqrt` from the `math` module is fast. Note that there seems to be no difference between `math.sqrt` and `sqrt`, which is just a local reference. Let's look at this in a further example:

```
# local_ref.py  
  
"""Testing access to local name and name referenced in another module.  
"""  
  
import math  
import sys  
  
if sys.version_info.major < 3:  
    range = xrange  
  
# If there is no decorator `profile`, make one that just calls the function,  
# i.e. does nothing.  
# This allows to call `kernprof` with and without the option `-l` without
```

2.7 Going Line-by-Line

```
# commenting or un-commenting `@profile` all the time.
# You can add this to the builtins to make it available in the whole program.
try:
    @profile
    def dummy():
        """Needs to be here to avoid a syntax error.
        """
        pass
except NameError:
    def profile(func):
        """Will act as the decorator `profile` if it was already found.
        """
        return func

@profile
def local_ref(counter):
    """Access local name.
    """
    # make it local
    sqrt = math.sqrt
    for _ in range(counter):
        sqrt

@profile
def module_ref(counter):
    """Access name as attribute of another module.
    """
    for _ in range(counter):
        math.sqrt

@profile
def test(counter):
    """Call both functions.
    """
    local_ref(counter)
    module_ref(counter)

if __name__ == '__main__':
    test(int(1e7))
```

There are two functions to be line-traced. `local_ref` gets a local reference to `math.sqrt` and `module_ref` calls `math.sqrt` as it is.

We run this with the option `-v`, and we get:

```
$ kernprof -v local_ref.py
Wrote profile results to local_ref.py.prof
 9 function calls in 14.847 seconds
```

2.7 Going Line-by-Line

```
Ordered by: standard name

ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
      1     0.000      0.000    14.847    14.847 <string>:1(<module>)
      1     0.000      0.000     0.000     0.000 local_ref.py:18(profile)
      1     0.001      0.001    14.846    14.846 local_ref.py:2(<module>)
      1     0.000      0.000    14.845    14.845 local_ref.py:21(mock)
      1     4.752      4.752     4.752     4.752 local_ref.py:28(local_ref)
      1    10.093    10.093    10.093    10.093 local_ref.py:37(module_ref)
      1     0.000      0.000    14.845    14.845 local_ref.py:44(test)
      1     0.001      0.001    14.847    14.847 {execfile}
      1     0.000      0.000     0.000     0.000 {method 'disable' of
'_lsprof.Profiler' objects}
```

This shows that `local_ref` is more than twice as fast as `module_ref` because it avoids many lookups on the module `math`.

Now we run it with the options `-v -l`:

```
$ kernprof -v -l local_ref.py
Wrote profile results to local_ref.py.lprof
Timer unit: 1e-06 s

File: local_ref.py
Function: dummy at line 12
Total time: 0 s

Line #      Hits          Time  Per Hit   % Time  Line Contents
=====
     12                               @profile
     13                               def dummy():
     14                                   """Needs to be here to
avoid a syntax error.
     15                                   """
     16                                   pass

File: local_ref.py
Function: test at line 44
Total time: 125.934 s

Line #      Hits          Time  Per Hit   % Time  Line Contents
=====
     44                               @profile
     45                               def test(counter):
     46                                   """Call both functions.
     47                                   """
     48                               1    58162627 58162627.0    46.2    local_ref(counter)
     49                               1    67771433 67771433.0    53.8    module_ref(counter)
```

2.7 Going Line-by-Line

This takes much longer. The differences in run times are largely gone. After correspondence with Robert Kern, the author of `line_profiler`, it turns out that the substantial overhead the line tracing adds causes a distortion of measuring results. Conclusion: Use `line_profiler` for expensive atomic calls such as to a function in an extension module like NumPy.

In the Jupyter Notebook, you can also use the `%lprun` magic command. First, you need import the extension:

```
%load_ext line_profiler
```

Just use the `-f` flag instead of the decorator `profile`:

```
%lprun -f use_fast fast use_fast()
```

to get this output:

```
Timer unit: 1e-06 s

Total time: 0.12277 s
File: <ipython-input-11-481ad1301326>
Function: fast at line 12

Line #      Hits          Time  Per Hit   % Time  Line Contents
=====
    12                               def fast():
    13                               """Wait 0.001 seconds.
    14                               """
    15          100         122770   1227.7    100.0    time.sleep(1e-3)

Total time: 0.123318 s
File: <ipython-input-11-481ad1301326>
Function: use_fast at line 24

Line #      Hits          Time  Per Hit   % Time  Line Contents
=====
    24                               def use_fast():
    25                               """Call `fast` 100 times.
    26                               """
    27          101          146      1.4      0.1    for _ in range(100):
    28          100        123172  1231.7    99.9        fast()
```

2.7.1 Exercise

Profile these two functions with `cProfile` and with `%prun` in an Jupyter Notebook as well as with the line profiler. Save the results and visualize them with SnakeVis.

```
# file: create_list.py

import sys
```



```

if sys.version_info.major < 3:
    range = xrange

def insert_zero(n=int(1e4)):
    """Assemble list with `insert`. Inefficient.
    """
    L = []
    for x in range(n):
        L.insert(0, x)
    return L

def append_reverse(n=int(1e4)):
    """Assemble list with `append` and `reverse`.
    """
    L = []
    for x in range(n):
        L.append(x)
    L.reverse()
    return L

```

2.8 Profiling Memory Usage

Current computers have lots of RAM, still it can be a problem if an application uses more RAM than is physically available, leading to swapping and a large performance penalty. In particular, long running applications tend to use up more RAM over time. Although Python does automatic memory management, there are cases where memory is not released because there are still references to objects that are no longer needed. We can call the garbage collector manually, but this does not always produce the desired effects.

2.8.1 Pympler

[Pympler](#)³ it is a merge of the formerly independent projects `asizeof`, `heapmonitor`, and `muppy`.

Let's start a new interpreter and make an instance of `pympler.tracker.SummaryTracker`:

```

>>> from pympler import tracker
>>> mem_tracker = tracker.SummaryTracker()

```

We need to call `print_diff()` several times to get to the baseline:

```

>>> mem_tracker.print_diff()

```

types	# objects	total size
list	1353	138.02 KB
str	1345	75.99 KB
int	149	3.49 KB
dict	2	2.05 KB

2.8 Profiling Memory Usage

```
wrapper_descriptor |          8 |        640 | B
weakref            |          3 |        264 | B
member_descriptor |          2 |        144 | B
getset_descriptor |          2 |        144 | B
function (store_info) |        1 |        120 | B
cell              |          2 |        112 | B
instancemethod    |         -1 |         -80 | B
tuple            |         -1 |       -216 | B
>>> mem_tracker.print_diff()
types | # objects | total size
=====|=====|=====
str   |          2 |         97 | B
list  |          1 |         96 | B
>>> mem_tracker.print_diff()
types | # objects | total size
=====|=====|=====
```

Now we create our big list and look at the memory again:

```
>>> big_list = list(range(int(1e6)))
>>> mem_tracker.print_diff()
types | # objects | total size
=====|=====|=====
int   |    999861 |    22.89 MB
list  |          1 |     7.63 MB
```

Let's look at some examples for how we can use `pympler`. First we write a decorator that tells us how much memory the result of a function uses:

```
# file: memory_size_pympler.py

"""Measure the size of used memory with a decorator.
"""

from __future__ import print_function

import functools #1
import sys

if sys.version_info.major < 3:
    range = xrange

from pympler import tracker #2

memory = {} #3

def measure_memory(function): #4
```

2.8 Profiling Memory Usage

```
"""Decorator to measure memory size.
"""

@functools.wraps(function) #5
def _measure_memory(*args, **kwargs): #6
    """This replaces the function that is to be measured.
    """
    measurer = tracker.SummaryTracker() #7
    for _ in range(2): #8
        measurer.diff() #9
    try:
        res = function(*args, **kwargs) #10
        return res
    finally: #11
        memory[function.__name__] = (measurer.diff())
    return _measure_memory #12

if __name__ == '__main__':

    @measure_memory #13
    def make_big(number):
        """Example function that makes a large list.
        """
        return list(range(number)) #14

    make_big(int(1e6)) #15
    print('used memory', memory) #16
```

First we import `functools` (#1) that will help us to write a nice decorator. Then we import `pympler.tracker` (#2) and define a global dictionary (#3) that will hold all values for memory. We define a function that takes a function as argument (#4) and another function inside it that takes a variable number of positional and keyword arguments (#6). This is a typical setup of a decorator that takes no arguments (with arguments we would need a third level). We also decorate this function with `@functools.wraps` (#5) to preserve the docstring and the name of the original function after it is decorated.

Now we make an instance of our tracker (#7). We use a loop (#8) and call `tracker.diff()` several times (#9). Then we call the function with the supplied arguments (#10). We always want to have the size of memory after the call (#11). Finally, we return our internally defined function. Note that we store the result of the called function in `res`. This is necessary to get the memory that is used by the object the function returns. We return our newly created function (#12).

We decorate our function (#13) that just returns a list of size `number` (#14). After we call the function (#15), we can print the used memory (#16).

When we suspect that a function leaks memory, we can use `pympler` to measure the memory growth after a function returned:

```
# file memory_growth_pympler.py
```

2.8 Profiling Memory Usage

```
"""Measure the memory growth during a function call.
"""
from __future__ import print_function

import sys

if sys.version_info.major < 3:
    range = xrange

from pympler import tracker #1

def check_memory_growth(function, *args, **kwargs): #2
    """Measure the memory usage of `function`.
    """
    measurer = tracker.SummaryTracker() #3
    for _ in range(2): #4
        measurer.diff() #5
    function(*args, **kwargs) #6
    return measurer.diff() #7

if __name__ == '__main__':

    def test():
        """Do some tests with different memory usage patterns.
        """

        def make_big(number): #8
            """Function without side effects.

            It cleans up all used memory after it returns.
            """
            return list(range(number))

        data = [] #9

        def grow(number):
            """Function with side effects on global list.
            """
            for x in range(number):
                data.append(x) #10
            size = int(1e6)
            print('memory make_big:', check_memory_growth(make_big, #11
                                                            size))
            print('memory grow:', check_memory_growth(grow, size)) #12

        test()
```

2.8 Profiling Memory Usage

After importing `pympler.tracker` (#1) we define a helper function that takes the function to be measured, and positional and keyword arguments that will be handed to this function (#2). We make an instance of `tracker.SummaryTracker` (3) and use a loop (#4) to call `measurer.diff()` several times. This way, we set the baseline of memory usage (#5). We call the function with the supplied arguments (#6). Finally, we return the difference in memory size before and after the function call (#7).

We define a function that just returns a list (#8) and thus does not increase memory size after it is finished. The size of the returned list is not measured.

We use a global list as data storage (#9) and define a second function that appends elements to this list (#10). Finally, we call our helper function with both functions as arguments (#11 and #12).

Pympler offers more tools. Let's look at the possibilities to measure the memory size of a given object. We would like to measure the memory size of a list as we append elements. We write a function that takes the length of the list and a function that is to be used to measure the memory size of an object:

```
# file: pympler_list_growth.py

"""Measure the size of a list as it grows.
"""
from __future__ import print_function

import sys

from pympler.asizeof import asizeof, flatsize

if sys.version_info.major < 3:
    range = xrange

def list_mem(length, size_func=flatsize):
    """Measure incremental memory increase of a growing list.
    """
    my_list = []
    mem = [size_func(my_list)]
    for elem in range(length):
        my_list.append(elem)
        mem.append(size_func(my_list))
    return mem
```

Now we use this function with three different functions: `pympler.asizeof.flatsize`, `pympler.asizeof.asizeof` and `sys.getsizeof`:

```
if __name__ == '__main__':

    def main():
        """Show plot or numbers.
        """
        SIZE = 1000
```


2.8 Profiling Memory Usage

```
SHOW = 20

for func in [flatsize, asizeof, sys.getsizeof]:
    mem = list_mem(SIZE, size_func=func)
    try:
        from matplotlib import pylab
        pylab.plot(mem)
        pylab.show()
    except ImportError:
        print('matplotlib seems not be installed. Skipping the plot.')
    if SIZE > SHOW:
        limit = SHOW // 2
        print(mem[:limit],
              '... skipping %d elements ...' % (SIZE - SHOW),
              end='')
        print(mem[-limit:])
    else:
        print(mem)
```


The code just calls our function and supplies one of the functions to measure memory size as an argument. If matplotlib is installed, it draws a graph for each call. Let's look at the resulting graphs.

Using `pympler.asizeof.flatsize` we get this kind of step diagram:



We can see nicely how the list grows discontinuously. Python allocates more memory than it actually needs to append the next element. This way it can append several elements before it needs to increase its size again. These steps get bigger the bigger the list is.


Using `pympler.asizeof.asizeof` we get a different looking graph:



This function also measures the size of all referenced objects. In our case all the integer that are stored in the list. Therefore, there is an continuous increase in memory size between the steps case by the list allocation.

2.8 Profiling Memory Usage

For this simple case `sys.getsizeof` produces the same result as `pympler.asizeof.flatsize`:



For more complex cases `pympler.asizeof.flatsize` might give different results.

We can also measure the number of allocation steps it takes when a list grows one element at a time:

```
# file: list_alloc_steps.py

"""Measure the number of memory allocation steps for a list.
"""
from __future__ import print_function

import sys

if sys.version_info.major < 3:
    range = xrange

from pympler.asizeof import flatsize

def list_steps(lenght, size_func=sys.getsizeof):
    """Measure the number of memory alloaction steps for a list.
    """
    my_list = []
    steps = 0
    int_size = size_func(int())
    old_size = size_func(my_list)
    for elem in range(lenght):
        my_list.append(elem)
        new_size = sys.getsizeof(my_list)
        if new_size - old_size > int_size:
            steps += 1
        old_size = new_size
    return steps

if __name__ == '__main__':
    steps = [10, 100, 1000, 10000, int(1e5), int(1e6), int(1e7)]
    print('Using sys.getsizeof:')
    for size in steps:
        print('%10d: %3d' % (size, list_steps(size)))
    print('Using pympler.asizeof.flatsize:')
```

2.8 Profiling Memory Usage

```
for size in steps:
    print('%10d: %3d' % (size, list_steps(size, flatsize)))
```

The results are the same for `sys.getsizeof` and `pympler.asizeof.flatsize`:

```
Using sys.getsizeof:
    10: 3
   100: 10
  1000: 27
 10000: 46
100000: 65
1000000: 85
10000000: 104
Using pympler.asizeof.flatsize:
    10: 3
   100: 10
  1000: 27
 10000: 46
100000: 65
1000000: 85
10000000: 104
```

2.8.2 Memory Usage Line-by-Line with `memory_profiler`

Similarly to `line_profiler` that profiles CPU usage line-by-line, `memory_profiler` measures the memory line-by-line. We use a small sample code with one function and decorate it with `@profile`:

```
# file: use_mem.py

import random
import sys

# Make it work with Python 2 and Python 3.
if sys.version_info.major < 3:
    range = xrange

@profile
def use_mem(numbers):
    """Different ways to use up memory.
    """
    a = sum([x * x for x in numbers])
    b = sum(x * x for x in numbers)
    c = sum(x * x for x in numbers)
    squares = [x * x for x in numbers]
    d = sum(squares)
    del squares
    x = 'a' * int(1e6)
```


2.8 Profiling Memory Usage

```
del x
return 42

if __name__ == '__main__':
    numbers = [random.random() for x in range(int(1e6))]
    use_mem(numbers)
```

Running it from the command line:

```
$ python -m memory_profiler use_mem.py
```

for a list one million random numbers:

Line #	Mem usage	Increment	Line Contents
8			@profile
9	33.430 MB	0.000 MB	def use_mem(numbers):
10	94.797 MB	61.367 MB	a = sum([x * x for x in numbers])
11	94.797 MB	0.000 MB	b = sum(x * x for x in numbers)
12	94.797 MB	0.000 MB	c = sum(x * x for x in numbers)
13	114.730 MB	19.934 MB	squares = [x * x for x in numbers]
14	121.281 MB	6.551 MB	d = sum(squares)
15	121.281 MB	0.000 MB	del squares
16	312.020 MB	190.738 MB	x = 'a' * int(2e8)
17	121.281 MB	-190.738 MB	del x
18	121.281 MB	0.000 MB	return 42

and then for a list ten million random numbers:

Line #	Mem usage	Increment	Line Contents
8			@profile
9	265.121 MB	0.000 MB	def use_mem(numbers):
10	709.500 MB	444.379 MB	a = sum([x * x for x in numbers])
11	799.570 MB	90.070 MB	b = sum(x * x for x in numbers)
12	798.965 MB	-0.605 MB	c = sum(x * x for x in numbers)
13	806.707 MB	7.742 MB	squares = [x * x for x in numbers]
14	972.270 MB	165.562 MB	d = sum(squares)
15	976.984 MB	4.715 MB	del squares
16	943.906 MB	-33.078 MB	x = 'a' * int(2e8)
17	871.207 MB	-72.699 MB	del x
18	871.203 MB	-0.004 MB	return 42

The result is not as clear as expected. One reason might be that it takes time to free memory. Therefore, the effects come later.

2.8 Profiling Memory Usage

In addition to running from the command line you can import the decorator from `memory_profile` import `profile`. You can also track the memory usage over time. For example, `his` measures the usage of the interactive Python interpreter:

```
>>> from memory_profiler import memory_usage
>>> mem_over_time = memory_usage(-1, interval=0.5, timeout=3)
>>> mem_over_time
[7.453125, 7.4609375, 7.4609375, 7.4609375, 7.4609375, 7.4609375]
```

You can also supply a PID of another process. `memory_profiler` also comes with a IPython plug-in to be used with the magic function `%memit` analogous to `%timeit`. You need to enable it with `%load_ext memory_profiler`.

We can also use our line-by-line memory profiler in a Jupyter Notebook.

First, we need to load the extension:

```
%load_ext memory_profiler
```

Now we can import our function:

```
from use_mem_no_deco import use_mem
```

and use the magic command `%mprun`, indicating with `-f` what function should be measured:

```
%mprun -f use_mem use_mem([random.random() for x in range(int(1e6))])
```

The result looks like this:

```
Filename: ../use_mem_no_deco.py
```

Line #	Mem usage	Increment	Line Contents
11	65.1 MiB	0.0 MiB	def use_mem(numbers):
12			"""Different ways to use up memory.
13			"""
14	96.7 MiB	31.6 MiB	a = sum([x * x for x in numbers])
15	75.7 MiB	-21.0 MiB	b = sum(x * x for x in numbers)
16	75.7 MiB	0.0 MiB	c = sum(x * x for x in numbers)
17	98.1 MiB	22.5 MiB	squares = [x * x for x in numbers]
18	103.3 MiB	5.1 MiB	d = sum(squares)
19	76.2 MiB	-27.1 MiB	del squares
20	77.1 MiB	1.0 MiB	x = 'a' * int(1e6)
21	76.9 MiB	-0.2 MiB	del x
22	76.9 MiB	0.0 MiB	return 42

2.8.3 Exercise

Measure the memory consumption of these functions `make_list()` and `make_gen()` with Pympler. Measure the memory consumption of `test()` line-by-line with `memory_profiler`. Use increasing numbers for `n` from `1e4` to `1e7` (or larger if you have enough memory). Explain your findings.

```
#file: make_list_gen.py

from __future__ import print_function

import sys
import time

if sys.version_info.major < 3:
    range = xrange

def make_list(n):
    return [x * 10 for x in range(n)]

def make_gen(n):
    return (x * 10 for x in range(n))

def test():
    n = int(1e4) # 1e5, 1e6, 1e7, 1e8
    list_ = make_list(n)
    del list_
    gen = make_gen(n)
    del gen
    time.sleep(1)

test()
```

3 Algorithms and Anti-patterns

3.1 String Concatenation

Strings in Python are immutable. So if you want to modify a string, you have to actually create a new one and use parts of the old one:

```
>>> s = 'old text'
>>> 'new' + s[-5:]
'new text'
```

This means that new memory has to be allocated for the string. This is no problem for a few hundred or thousand strings, but if you have to deal with millions of strings, memory allocation time may be considerably longer. The solution in Python is to use a list to hold the sub strings and join them with `''.join()` string method.

3.1.1 Exercise

Write a test program that constructs a very long string (containing up to one million characters). Use the idiom `s += 'text'` and the idiom `text_list.append('text')` plus `''.join(text_list)` in a function for each. Compare the two approaches in terms of execution speed.

Hint: You can use `timeit.default_timer()` to get the time since the last call to this function. Alternatively, you can use the module `timeit` or the function `measure_run_time` from the module `measure_time` in the directory `measuring`. If you have PyPy install, run the timings with it.

3.2 List and Generator Expressions

Python offers list comprehension as a short and very readable way to construct a list.

```
>>> import sys
>>> if sys.version_info.major < 3: # for Python 2 users
    range = xrange
>>> L = [x * x for x in range(10)]
>>> L
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
```

is a short form for:

```
>>> L = []
>>> for x in range(10):
...     L.append(x * x)
...
>>> L
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
```

If you are not interested in the list itself but rather some values computed from the whole list, you can use generator comprehension and avoid the list all together.

3.3 Think Global buy Local

```
>>> sum(x * x for x in range(10))
285
```

3.2.1 Exercise

Write a test program that calculates the sum of all squares of the numbers from zero to one million. Use the idiom `l.append` and list comprehension as well as a generator comprehension. Try it with `range` and `list(range)` in Python 3. If you are still using Python 2, try it with `range` and `xrange`. Use different numbers, e.g. smaller and bigger than one million.

Hint: You can use `timeit.default_timer()` to get the time since the last call to this function. Alternatively, you can use the magic command `%timeit` in the Notebook or the function `measure_run_time` which you can find in the file `measure_time.py` in the `measuring` directory.

3.3 Think Global buy Local

A great deal of things in Python are dynamic. This includes the lookup of variables. It follows the famous LGB local-global-built-in rule. If a variable name is not found in the local scope, Python looks for it in global and then in the built-in name space before raising an `NameError` when nothing was found.

Since every name space is a dictionary, it involves more look ups the more name spaces have to be searched. Therefore, local variables are faster than global variables. Let's look at an example:

```
# file: local_global.py

"""Local vs. built-in.
"""

import sys

if sys.version_info.major < 3:
    range = xrange

GLOBAL = 1

def repeat(counter):
    """Using the GLOBAL value directly.
    """
    for count in range(counter):
        GLOBAL

def repeat_local(counter):
    """Making GLOBAL a local variable.
    """
    local = GLOBAL
    for count in range(counter):
        local
```

3.3 Think Global buy Local

```
def test(counter):  
    """Call both functions.  
    """  
    repeat(counter)  
    repeat_local(counter)  
  
if __name__ == '__main__':  
  
    def do_profile():  
        """Check the run times.  
        """  
        import cProfile  
        profiler = cProfile.Profile()  
        profiler.run('test(int(1e8))')  
        profiler.print_stats()  
  
    do_profile()
```

By running this code, we will see that the version that accesses the GLOBAL directly is about 25% slower than the version with the local variable.

The difference becomes larger when we move more outward and make a built-in name a local one:

```
"""Local vs. built-in.  
"""  
  
import sys  
  
if sys.version_info.major < 3:  
    range = xrange  
  
def repeat(counter):  
    """Using the built-in `sum` in a loop.  
    """  
    for count in range(counter):  
        sum  
  
def repeat_local(counter):  
    """Making `sum` a local variable.  
    """  
    sum_ = sum  
    for count in range(counter):  
        sum_
```

3.3 Think Global buy Local

```
def test(counter):  
    """Call both functions.  
    """  
    repeat(counter)  
    repeat_local(counter)  
  
if __name__ == '__main__':  
  
    def do_profile():  
        """Check the run times.  
        """  
        import cProfile  
        profiler = cProfile.Profile()  
        profiler.run('test(int(1e8))')  
        profiler.print_stats()  
  
    do_profile()
```

In this example it saves about 40% of the run. So, if you have large loops and you access globals or built-ins frequently, making them local might be quite useful.

4 The Right Data Structure

4.1 Use built-in Data Types

It is always a good idea to use Python built-in data structures. They are not only most often more elegant and robust than self-made data structures, but also faster in nearly all cases. They are well tested, often partially implemented in C and optimized through long time usage by many talented programmers.

There are essential differences among built-in data types in terms of performance depending on the task.

4.2 list vs. set

If you need to search in items, dictionaries and sets are mostly preferable to lists.

```
>>> 9 in range(10)
True
>>> 9 in set(range(10))
True
```

Let's make a performance test. We define a function that searches in a list:

```
# file: searching.py
"""Measuring the time for searching in a list and a set.
"""

import timeit

def search_list(n):
    """
    Search for element that is not in a list.
    """
    my_list = list(range(n))
    start = timeit.default_timer()
    n in my_list # pylint: disable=pointless-statement
    return timeit.default_timer() - start
```

and one that searches in a set:

```
def search_set(n):
    """Search for an element in a set.
    """
    my_set = set(range(n))
    start = timeit.default_timer()
    n in my_set # pylint: disable=pointless-statement
    return timeit.default_timer() - start
```

We define a function that compares both run times:

4 The Right Data Structure

```
def calculate_ratio(n):  
    """Calculate the ratio between a search in a list and a set.  
    """  
    list_time = search_list(n)  
    set_time = search_set(n)  
    return list_time, set_time, list_time / set_time
```

as well as a helper function that creates a nice table for different sizes:

```
def compare(end=8, func=calculate_ratio, header='', coll='List', col2='Set'):  
    """Show the results.  
    """  
    table_width = 43  
    print()  
    if header:  
        print('=' * table_width)
```

Running this function:

```
if __name__ == '__main__':  
    compare(header='Single run')
```

shows that the set is considerably faster, especially for larger collections:

```
=====
Single run
=====
      Size      List      Set      Ratio
-----
      10  8.13e-07  2.59e-07      3.14
     100  1.34e-06  1.58e-07      8.51
    1,000  1.14e-05  2.05e-07     55.48
   10,000  1.14e-04  3.31e-07    346.06
  100,000  1.09e-03  6.92e-07   1,568.28
1,000,000  1.37e-02  1.44e-06   9,509.33
10,000,000 1.21e-01  1.84e-06  66,009.02
=====
```

The times for searching in the set should be the same for all sizes of the set. There might be some problem with our measurement. Let's write a function that measures similarly to IPython's `%timeit`, i.e. tries to run for about 1 second, repeats this seven times, and finally returns the average run time for this seven repeats. For example:

```
# file: searching_multiple.py  
"""  
Measuring the time for searching in a list and a set multiple times.  
"""
```

4 The Right Data Structure

```
from statistics import mean
import timeit

from searching import compare

def search_multiple(obj, n, repeat=7):
    """Search `repeat` times for at least 1 second.
    """
    res = []
    for _ in range(repeat):
        count = 0
        duration = 0
        while duration < 1:
            start = timeit.default_timer()
            n in obj # pylint: disable=pointless-statement
            duration += timeit.default_timer() - start
            count += 1
        res.append(duration / count)
    return mean(res)
```

We need a modified compare function:

```
def calculate_ratio_mutable(n):
    """Calculate the ratio between a search in a list and a set.
    """
    my_list = list(range(n))
    my_set = set(range(n))
    list_time = search_multiple(my_list, n)
    set_time = search_multiple(my_set, n)
    return list_time, set_time, list_time / set_time
```

Running this function:

```
if __name__ == '__main__':
    compare(func=calculate_ratio_mutable, header='Multiple runs')
```

we get different run times for the set:

```
=====
Multiple runs
=====
      Size      List      Set      Ratio
-----
      10  2.55e-07  1.34e-07      1.90
     100  1.31e-06  1.34e-07      9.80
```

4 The Right Data Structure

1,000	1.19e-05	1.33e-07	89.27
10,000	1.27e-04	1.33e-07	951.72
100,000	1.56e-03	1.34e-07	11,616.24
1,000,000	1.22e-02	1.34e-07	90,527.51
10,000,000	1.25e-01	1.38e-07	908,220.24

=====

Now, the set run times are the same for all sizes of the set. Starting from a size of 100, the list run times are comparable. This shows that our measurement error is large for short run times.

Let's explore this a bit further and use IPython's magic function `%timeit` (as used in the Jupyter Notebook) for our measurement. We can use the method `TerminalInteractiveShell().run_cell_magic()` to run it from a script instead of interactively from the Jupyter Notebook or the IPython shell. We write a small helper function:

```
# file: searching_magic.py
"""
Measuring the time for searching in a list and a set with IPython %timeit.
"""

from IPython.terminal.interactiveshell import TerminalInteractiveShell

from searching import compare

def timeit_magic(n, setup, statement):
    """Create a `%timeit` magic function with fixed `n`,
    more setup code and the statement to be timed.
    """
    return TerminalInteractiveShell().run_cell_magic(
        'timeit', '-o -q n = {n}'.format(n=n) + setup, statement)
```

and two timing functions.

One for a list:

```
def search_list(n):
    """
    Search for last element in a list.
    """
    setup = 'my_list = list(range(n))'
    statement = 'n in my_list'
    return timeit_magic(n, setup, statement)
```

and one for a set:

```
def search_set(n):
    """Search for an element in a set.
    """
```

4 The Right Data Structure

```
setup = 'my_set = set(range(n))'
statement = 'n in my_set'
return timeit_magic(n, setup, statement)
```

We also need a new ratio calculation function:

```
def calculate_ratio(n, search_list=search_list, search_set=search_set):
    """Calculate the ratio between a search in a list and a set.
    """
    list_time = search_list(n).average
    set_time = search_set(n).average
    return list_time, set_time, list_time / set_time
```

Finally, we run our function:

```
if __name__ == '__main__':
    compare(func=calculate_ratio, header='Magic timeit')
```

with this result:

```
=====
Magic timeit
=====
      Size      List      Set      Ratio
-----
      10  1.45e-07  2.99e-08      4.84
     100  1.20e-06  3.09e-08     38.81
    1,000  1.19e-05  3.14e-08    379.36
   10,000  1.18e-04  2.96e-08   3,990.61
  100,000  1.21e-03  3.10e-08  39,006.49
 1,000,000  1.19e-02  2.99e-08 398,732.10
10,000,000  1.24e-01  3.10e-08 4,000,213.05
=====
```

The timings for the list are essentially the same as for own timing functions. But the timings for the set are consistently about four times faster. Looks like that our measurements in `searching_multiple.py` are not very accurate for short run times.

We did not measure the time it takes to convert the list into a set. So, let's define a modified function for the set that includes the creation of the set from the list into the runtime measurement:

```
# file: searching_creation.py

"""
Measuring the time for searching in a list and a set including
creation time of the data structure.
"""
```

4 The Right Data Structure

```
from functools import partial
import textwrap

from searching import compare
from searching_magic import calculate_ratio, timeit_magic

def search_set(n):
    """Search for an element in a set.
    """
    setup = 'my_list = list(range(n))'
    statement = textwrap.dedent("""
    my_set = set(my_list)
    n in my_set
    """)
    return timeit_magic(n, setup, statement)

def main():
    """Run some timings.
    """
    func = partial(calculate_ratio, search_set=search_set)
    compare(func=func, header='Measure creation')

if __name__ == '__main__':
    main()
```

Now the set is not faster anymore:

```
=====
Measure creation
=====
      Size      List      Set      Ratio
-----
      10  1.52e-07  3.39e-07    0.45
     100  1.21e-06  1.95e-06    0.62
    1,000  1.19e-05  1.33e-05    0.90
   10,000  1.19e-04  1.58e-04    0.75
  100,000  1.19e-03  1.82e-03    0.65
1,000,000  1.20e-02  3.63e-02    0.33
10,000,000 1.23e-01  4.66e-01    0.26
=====
```

If we need to search more than once, the overhead for creating the set gets relatively smaller. We write a function that searches in our list several times:

4 The Right Data Structure

```
# file: searching_repeated.py
"""
Measuring the time for searching in a list and a set including
creation time of the data structure.
"""

from functools import partial
import textwrap

from searching import compare
from searching_magic import calculate_ratio, timeit_magic

def search_list(n, m):
    """Search for an element in a set.
    """
    setup = 'my_list = list(range(n));m = {m}'.format(m=m)
    statement = textwrap.dedent("""
    for x in range(m):
        n in my_list
    """)
    return timeit_magic(n, setup, statement)
```

and do the same for our set:

```
def search_set(n, m):
    """Search for an element in a set.
    """
    setup = 'my_list = list(range(n));m = {m}'.format(m=m)
    statement = textwrap.dedent("""
    my_set = set(my_list)
    for x in range(m):
        n in my_set
    """)
    return timeit_magic(n, setup, statement)
```

We start our measurement for different repeats:

```
if __name__ == '__main__':
    main()
```

This may take a while.

This will run for several minutes.

4 The Right Data Structure

If you get an error `OSError: [Errno 24] Too many open files`, try running the script twice, ones with `for m in [10, 100]:` and ones with `for m in [1000, 10000]:` to reduce the number of runs in a single Python process, or increase the number allowed open files (with a method depending on your operating system).

The set gets relatively faster with increasing collection size and number of searches:

```
=====
Measure for 10 repetitions
=====
      Size      List      Set      Ratio
-----
      10  1.75e-06  9.29e-07      1.88
     100  1.23e-05  2.52e-06      4.90
    1,000  1.18e-04  1.39e-05      8.44
   10,000  1.18e-03  1.57e-04      7.52
  100,000  1.18e-02  1.81e-03      6.54
1,000,000  1.20e-01  3.87e-02      3.09
10,000,000 1.31e+00  4.72e-01      2.77
=====

=====
Measure for 100 repetitions
=====
      Size      List      Set      Ratio
-----
      10  1.49e-05  3.99e-06      3.74
     100  1.19e-04  5.58e-06     21.35
    1,000  1.18e-03  1.69e-05     69.88
   10,000  1.17e-02  1.60e-04     73.06
  100,000  1.19e-01  1.95e-03     61.11
1,000,000  1.19e+00  3.36e-02     35.38
10,000,000 1.21e+01  4.76e-01     25.52
=====

=====
Measure for 1000 repetitions
=====
      Size      List      Set      Ratio
-----
      10  1.53e-04  4.17e-05      3.67
     100  1.20e-03  4.65e-05     25.71
    1,000  1.19e-02  5.89e-05    201.84
   10,000  1.19e-01  2.07e-04    577.73
  100,000  1.20e+00  1.88e-03    637.18
1,000,000  1.20e+01  3.67e-02    327.16
10,000,000 1.22e+02  4.57e-01    268.32
```

4 The Right Data Structure

```
=====
=====
Measure for 10000 repetitions
=====
      Size      List      Set      Ratio
-----
      10    1.59e-03    4.52e-04      3.51
     100    1.23e-02    4.58e-04     26.78
    1,000    1.14e-01    4.63e-04    246.34
   10,000    1.18e+00    6.30e-04   1,869.65
  100,000    1.19e+01    2.34e-03   5,071.73
1,000,000    1.20e+02    3.47e-02   3,461.97
10,000,000    1.21e+03    4.70e-01   2,566.88
=====
```

Let's assume we have two lists:

```
>>> list_a = list('abcdefg')
>>> list_a
['a', 'b', 'c', 'd', 'e', 'f', 'g']
>>> list_b = list('fghijklmnopq')
>>> list_b
['f', 'g', 'h', 'i', 'j', 'k', 'l', 'm', 'n', 'o', 'p', 'q']
```

and we would like to find out which letters are in both lists. A simple implementation would look like this:

```
>>> in_both = []
>>> for a in list_a:
...     if a in list_b:
...         in_both.append(a)
```

```
>>> in_both
['f', 'g']
```

This can be achieved in fewer lines and in most cases faster with sets:

```
>>> set_a = set(list_a)
>>> set_b = set(list_b)
>>> set_a.intersection(set_b)
set(['g', 'f'])
```

Following the same method, we write a short performance test. First we write the function that uses lists:

```
# file: intersect.py
"""
Measuring the time for searching in a list and a set including
```


4 The Right Data Structure

```
creation time of the data structure.
"""

import timeit

from searching import compare

def intersect_list(n):
    """Measure the run time for intersecting two lists.
    """
    list_a = range(n)
    list_b = range(n-3, 2 * n)
    start = timeit.default_timer()
    in_both = []
    for x in list_a:
        if x in list_b:
            in_both.append(x)
    run_time = timeit.default_timer() - start
    return run_time, in_both
```

Now, we write a function for sets:

```
def intersect_set(n):
    """Measure the run time for intersecting two setss.
    """
    set_a = set(range(n))
    set_b = set(range(n-3, 2 * n))
    start = timeit.default_timer()
    in_both = set_a.intersection(set_b)
    run_time = timeit.default_timer() - start
    return run_time, in_both
```

We also write a function that calls both intersection functions and checks that both return the same result. It also returns the run times:

```
def calculate_intersect(n):
    """Calculate the intersecting time for two lists and two sets.
    """
    list_time, list_result = intersect_list(n)
    set_time, set_result = intersect_set(n)
    assert set_result == set(list_result)
    return list_time, set_time, list_time / set_time
```

Finally, we run out comparison:

4 The Right Data Structure

```
if __name__ == '__main__':  
    compare(func=calculate_intersect, header='Intersection')
```

with this result:

```
=====
Intersection
=====
      Size      List      Set      Ratio
-----
      10  5.72e-06  1.13e-06      5.07
     100  6.53e-06  1.12e-06      5.82
    1,000  8.45e-05  7.94e-06     10.65
   10,000  6.81e-04  7.89e-05      8.64
  100,000  7.35e-03  7.42e-04      9.91
1,000,000  7.37e-02  8.70e-03      8.48
10,000,000 7.23e-01  4.40e-01      1.64
=====
```

We do our run only ones. Let's get a second opinion and implement a another test script that uses IPython's `%timeit`.

We define function to measure the run time for intersection two lists:

```
# file: intersect_magic.py
"""
Measuring the time for searching in a list and a set including
creation time of the data structure with IPython's `%timeit`.
"""

from functools import partial
import textwrap

from searching import compare
from searching_magic import calculate_ratio, timeit_magic

def intersect_list(n):
    """Measure the run time for intersecting two lists.
    """
    setup = 'list_a = range(n);list_b = range(n-3, 2 * n)'
    statement = textwrap.dedent("""
in_both = []
for a in list_a:
    if a in list_b:
        in_both.append(a)
""")
    return timeit_magic(n, setup, statement)
```

4.3 list vs. deque

and one for two sets:

```
def intersect_set(n):  
    """Measure the run time for intersecting two sets.  
    """  
    setup = 'set_a = set(range(n));set_b = set(range(n-3, 2 * n))'  
    statement = 'set_a.intersection(set_b)'  
    return timeit_magic(n, setup, statement)
```

Now we measure:

```
if __name__ == '__main__':  
    main()
```

Our results differs a bit from the ones obtain with our single-timing approach above:

```
=====
Intersection magic
=====
      Size      List      Set      Ratio
-----
      10    1.02e-06    2.08e-07         4.90
     100    6.73e-06    9.03e-07         7.45
    1,000    7.16e-05    5.76e-06        12.42
   10,000    7.38e-04    6.46e-05        11.41
  100,000    7.37e-03    6.26e-04        11.78
1,000,000    7.34e-02    6.17e-03        11.90
10,000,000    7.60e-01    3.69e-01         2.06
=====
```

Note that the time for constructing the sets is not included in these measurements.

4.3 list vs. deque

For certain tasks we can use a deque instead of a list. A deque is a doubly linked list. This data structure allows faster insertion into the middle part. On the other hand, access of elements by index is slow.

We would like to achieve this:

```
>>> L = range(10)
>>> L[2:4] = []
>>> L
[0, 1, 4, 5, 6, 7, 8, 9]
```

Let's try the same with a deque. We import deque from the collections module:

```
>>> from collections import deque
```

make a deque:

4.3 list vs. deque

```
>>> d = deque(range(10))
>>> d
deque([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

Now, we rotate by the negative end index:

```
>>> d.rotate(-4)
>>> d
deque([4, 5, 6, 7, 8, 9, 0, 1, 2, 3])
```

We remove the last two elements:

```
>>> d.pop()
3
>>> d.pop()
2
```

and rotate back in the desired order:

```
>>> d.rotate(2)
>>> d
deque([0, 1, 4, 5, 6, 7, 8, 9])
```

Now, let's test the speed. We write test function for the measuring speed.

```
# file: list_deque.py

"""Removing elements from a list vs. from a deque.
"""

from collections import deque
from statistics import mean
import timeit

def time_function(func, make_args, repeat=7, limit=1):
    """Measure the run time of a function."""
    timing_res = []
    for _ in range(repeat):
        count = 0
        duration = 0
        while duration < limit:
            args = make_args()
            start = timeit.default_timer()
            func(*args)
            duration += timeit.default_timer() - start
            count += 1
```

4.3 list vs. deque

```
    timing_res.append(duration / count)
    return mean(timing_res)
```

This not that easy. We need to re-create the structures for each test but also want to run the removal many times to get reliable measurements.

We define a function to replace elements of a list:

```
def remove_from_list(my_list, start, end):
    """Remove elements between `start` and `end` from a list.
    """
    my_list[start:end] = []
```

and one to replace elements of a deque:

```
def remove_from_deque(my_deque, start, end):
    """Remove elements between `start` and `end` from a deque.
    """
    my_deque.rotate(-end)
    for _ in range(end - start):
        my_deque.pop()
    my_deque.rotate(start)
```

The main function measures for different limits, i.e. times to run for each repeat. The re-creation of data structures is very expensive. Therefore this will take some time:

```
def main():
    """Run some tests.
    """
    start = 100
    size = int(1e6)
    fmt = '{diff:10,d} {list_time:10.2e} {deque_time: 10.2e} {ratio:8.2f}'
    for limit in [0.00001, 0.0001, 0.001, 0.01, 0.1]:
        print('Limit:', limit)
        print('{:>10s} {:>10s} {:>10s} {:>8s}'.format(
            'Replaced', 'List', 'Deque', 'Ratio'))
        for end in [101, 110, 1100, 10100, 100100]:
            diff = end - start
            results = {}
            for obj, func in zip([list, deque], [remove_from_list,
                                                remove_from_deque]):
                def make_args(obj=obj, size=size, start=start, end=end):
                    """Dynamically create function with right arguments.
                    """
                    return obj(range(size)), start, end

                res = time_function(func, make_args, limit=limit)
                results[obj.__name__] = res
```

4.3 list vs. deque

```
list_time = results['list']
deque_time = results['deque']
ratio = list_time / deque_time
print(fmt.format(diff=diff, list_time=list_time,
                 deque_time=deque_time, ratio=ratio))

if __name__ == '__main__':
    main()
```

The results:

```
Limit: 1e-05
Replaced    List      Deque      Ratio
    1    6.53e-04    1.09e-05    59.73
   10    6.91e-04    1.13e-05    60.99
  1,000    6.69e-04    1.20e-04     5.59
 10,000    8.51e-04    8.50e-04     1.00
100,000    2.45e-03    8.48e-03     0.29

Limit: 0.0001
Replaced    List      Deque      Ratio
    1    7.30e-04    1.06e-05    69.15
   10    7.14e-04    1.16e-05    61.80
  1,000    8.40e-04    9.75e-05     8.62
 10,000    8.67e-04    8.86e-04     0.98
100,000    2.33e-03    8.72e-03     0.27

Limit: 0.001
Replaced    List      Deque      Ratio
    1    8.01e-04    1.07e-05    75.05
   10    6.98e-04    1.04e-05    67.30
  1,000    7.27e-04    9.84e-05     7.38
 10,000    8.81e-04    8.77e-04     1.00
100,000    2.39e-03    8.85e-03     0.27

Limit: 0.01
Replaced    List      Deque      Ratio
    1    7.73e-04    1.03e-05    74.74
   10    6.88e-04    1.09e-05    63.05
  1,000    7.03e-04    1.02e-04     6.92
 10,000    9.23e-04    9.28e-04     0.99
100,000    2.55e-03    8.87e-03     0.29

Limit: 0.1
Replaced    List      Deque      Ratio
    1    6.93e-04    1.08e-05    63.94
   10    6.38e-04    1.18e-05    54.28
  1,000    6.65e-04    9.96e-05     6.68
 10,000    1.02e-03    9.52e-04     1.07
100,000    2.68e-03    9.58e-03     0.28
```

reveals that we get reasonable measurements also for relative small limits.

4.4 dict vs. defaultdict

There is a `defaultdict` in the module `collections`. This works similarly to the `setdefault` method of ordinary dictionaries.

Let's assume we want to count how many of each letter are in the following sentence:

```
>>> s = 'Some letters appear several times in this text.'
```

We can do this in the standard way:

```
>>> d = {}
>>> for key in s:
...     d.setdefault(key, 0)
...     d[key] += 1
...
>>> d
{'a': 3, ' ': 7, 'e': 8, 'i': 3, 's': 4, 'm': 2,
 'l': 2, 'o': 1, 'n': 1, 'p': 2, 'S': 1, 'r': 3,
 't': 6, 'v': 1, 'x': 1, 'h': 1, '.': 1}
```

Or we can use the new `defaultdict`:

```
>>> dd = collections.defaultdict(int)
>>> for key in s:
...     dd[key] += 1
...
>>> dd
defaultdict(<type 'int'>, {'a': 3, ' ': 7, 'e': 8, 'i': 3, 's': 4, 'm': 2,
 'l': 2, 'o': 1, 'n': 1, 'p': 2, 'S': 1, 'r': 3, 't': 6, 'v': 1, 'x': 1,
 'h': 1, '.': 1})
>>>
```

Let's profile the speed differences. First, a function with our standard dictionary:

```
# file: setdefault_defaultdict.py

"""Defaultdict can faster than a standard dict.
"""

from collections import defaultdict

def standard_dict(text):
    """Count with standard dict.
    """
    d = {}
    for key in text:
        d.setdefault(key, 0)
```

4.4 dict vs. defaultdict

```
d[key] += 1
return d
```

And now one for the defaultdict:

```
def default_dict(text):
    """Count with defaultdict.
    """
    dd = defaultdict(int)
    for key in text:
        dd[key] += 1
    return dd
```

For a change, let's do our timing interactively in a Jupyter Notebook. We define our text:

```
text1 = 'Some letters appear several times in this text.'
```

Time our standard dictionary:

```
%timeit standard_dict(text1)
9.79 µs ± 90.3 ns per loop (mean ± std. dev. of 7 runs, 100000 loops each)
```

and our defaultdict:

```
%timeit default_dict(text1)
6.92 µs ± 81.8 ns per loop (mean ± std. dev. of 7 runs, 100000 loops each)
```

There is not much difference between them. Therefore, we increase the size of our data:

```
text2 = 'a' * 10 **6
```

Now the difference between standard dictionary:

```
%timeit standard_dict(text2)
201 ms ± 3.33 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

and defaultdict:

```
%timeit default_dict(text2)
91.4 ms ± 482 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)
```


4.4 dict vs. defaultdict

is more pronounced but not nearly in the region as our other speedups for list vs. sets or inserting a few elements in a large deque vs. doing this for a list.

There is an even better datastructure for our tasks, the Counter:

```
from collections import Counter
```

It is faster for the small text:

```
%timeit Counter(text1)
5.33 µs ± 63.4 ns per loop (mean ± std. dev. of 7 runs, 100000 loops each)
```

and even more so for the large text:

```
%timeit Counter(text2)
47.9 ms ± 569 µs per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

Let's look at a different example from the Python documentation. We have this data structure:

```
>>> data = [('yellow', 1), ('blue', 2), ('yellow', 3), ('blue', 4), ('red', 1)]
```

Our goal is to produce a dictionary that groups all second tuple entries into a list:

```
>>> d.items()
[('blue', [2, 4]), ('red', [1]), ('yellow', [1, 3])]
```

Again, we define a function for the dictionary versions:

```
def standard_dict_group(data):
    """Group with standard dict.
    """
    d = {}
    for key, value in data:
        d.setdefault(key, []).append(value)
    return d
```

and one for the defaultdict version:

```
def default_dict_group(data):
    """Group with defaultdict.
    """
    dd = defaultdict(list)
    for key, value in data:
```

4.5 Big-O notation and Data Structures

```
dd[key].append(value)
return dd
```

Now the standard dictionary is a bit faster:

```
%timeit standard_dict_group(data)
1.24 µs ± 14.8 ns per loop (mean ± std. dev. of 7 runs, 1000000 loops each)
```

than the defaultdict:

But, making the data much bigger:

```
data_large = data * 10000
```

This shows, that the standard dictionary needs about twice as long:

```
%timeit standard_dict_group(data_large)
9.35 ms ± 111 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)
```

as the defaultdict:

```
%timeit default_dict_group(data_large)
4.83 ms ± 48.1 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)
```

4.5 Big-O notation and Data Structures

Normally, you would want to reduce complexity of your program to make it faster. One frequently used measure for complexity is the so called **big-O** ⁴ notation. The following table gives an overview of some notations along with a short description and some examples from Python.

Notation	Description	Python Examples
$O(1)$	constant time does not increase with size of data	<code>len(my_list)</code> , <code>len(my_dict)</code> , <code>my_list[i]</code> , <code>del my_dict[i]</code> , <code>x in dict</code> , <code>x in set</code> , <code>my_list.append(i)</code>
$O(n)$	linear time increase linearly with size of data	Loops on list, strings, dicts, sets, string methods, <code>x in my_list</code>
$O(n \log n)$	quasi linear time increases a little faster than linearly	<code>my_list.sort()</code>
$O(n^2)$	quadratic time increases four times for each doubling of data	nested loops

4.6 $O(1)$ vs. $O(n)$ vs. $O(n^2)$

$O(n^3)$	cubic time increases eight times for each doubling of data	nested nested loops
$O(n!)$	factorial	traveling sales man problem (not Python specific)

In general, using big-O notation we look only at the order of magnitude. Constant factors are neglected. So $O(3*n)$ and $O(20*n)$ are called $O(n)$. Therefore, $O(20*n)$ might be slower than $O(n^2)$ for very small n . But for large n the constant factor has very little influence.

Actually we have already compared several of these notations in our examples above. Let's look at some more comparisons of notations.

4.6 $O(1)$ vs. $O(n)$ vs. $O(n^2)$

We write a function that takes an iterable and reverses it into a list. Our first implementation uses the method `insert` to insert every item at the first position:

```
from functools import partial
from statistics import mean
import timeit

from searching import compare
from searching_magic import calculate_ratio, timeit_magic

def use_on(iterable):
    result = []
    for item in iterable:
        result.insert(0, item)
    return result
```

Our second implementation uses `append` and reverses the list after all items are appended:

```
def use_o1(iterable):
    result = []
    for item in iterable:
        result.append(item)
    result.reverse()
    return result
```

We write a helper function to measure the run times:

```
def call_multiple(func, *args, repeat=7):
    """Search `repeat` times for at least 1 second.
    """
    res = []
    for _ in range(repeat):
        count = 0
        duration = 0
```

4.6 O(1) vs. O(n) vs. O(n²)

```
while duration < 1:
    start = timeit.default_timer()
    func(*args) # pylint: disable=pointless-statement
    duration += timeit.default_timer() - start
    count += 1
res.append(duration / count)
return mean(res)
```

Now we compare both functions in terms of runtime:

```
def main():
    """Run some timings.
    """
    print('{:>10s} {:>8s} {:>8s} {:>7s}'.format('Size', 'insert', 'append', 'Ratio'))
    fmt = '{size:10,d} {time_on:8.2e} {time_ol:8.2e} {ratio:7.2f}'
    for exp in range(1, 7):
        iterable = list(range(10 ** exp))
        time_on = call_multiple(use_on, iterable)
        time_ol = call_multiple(use_ol, iterable)
        print(fmt.format(size=len(iterable), time_on=time_on, time_ol=time_ol,
                        ratio=time_on/time_ol))

if __name__ == '__main__':
    main()
```

with these results:

Size	insert	append	Ratio
10	1.73e-06	1.11e-06	1.56
100	1.71e-05	7.79e-06	2.19
1,000	3.33e-04	7.19e-05	4.63
10,000	2.24e-02	7.01e-04	31.97
100,000	2.31e+00	8.04e-03	286.80
1,000,000	2.49e+02	8.21e-02	3038.70

The speed differences are growing rapidly with increasing data sizes. The method `append` is O(1) and `reverse` is O(n). Even though `insert` is also O(n) it is called n times whereas `reverse` is called only once.

Because we loop over all items of our iterable, the first function is O(n + n) but the second is O(n²). Of course instead of appending to a new list we can just convert the iterable into a list and reverse it:

```
def call_multiple(func, *args, repeat=7):
    """Search `repeat` times for at least 1 second.
    """
    res = []
    for _ in range(repeat):
```

4.7 Exercises

```
count = 0
duration = 0
while duration < 1:
    start = timeit.default_timer()
    func(*args) # pylint: disable=pointless-statement
    duration += timeit.default_timer() - start
    count += 1
res.append(duration / count)
return mean(res)
```

For a larger list:

```
iterable = list(range(10**6))
```

our $O(1)$ algorithm:

```
%timeit use_ol(iterable)
81.3 ms ± 2.32 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

is still about an order of magnitude slower, than the most efficient version:

Even though the big-O notation is the same, the `use_reversed` version is considerably faster.

4.7 Exercises

1. Modify the timing of searching in a list in `searching_multiple.py`. Instead of searching through the whole list, search for a target that is in the middle of the list. Hint: Use `n // 2` to find the middle index of the list. Run the timing and compare the results with the unmodified timing of searching in the list.

5 Caching

5.1 Reuse before You Recalculate

If you find yourself calling the same function with the same arguments many time then caching might help to improve the performance of your program. Instead of doing an expensive calculation, database query, or rendering again and over again, caching just reuses the results of former function calls. Depending on whether the results will be the same for every call to the same function with the same arguments or if the result might change over time, we talk about deterministic or non-deterministic caching. An example for deterministic caching would be numerical calculations that should always produce the same result for the same input. Caching of database queries is non-deterministic because the database content might change. So after some timeout period the query has to be done anew.

All of the following examples are based on [ZIAD2008].

5.2 Deterministic caching

The first thing we need to do, if we want to cache function results, is to uniquely identify the function we want to call:

```
# file: get_key.py
# based on Ziade 2008

"""Generate a unique key for a function and its arguments.
"""

def get_key(function, *args, **kw):                                #1
    """Make key from module and function names as well as arguments.
    """
    key = '%s.%s:' % (function.__module__,                        #2
                      function.__name__)
    hash_args = [str(arg) for arg in args]                        #3
    hash_kw = ['%s:%s' % (k, str(v))                             #4
               for k, v in kw.items()]
    return '%s::%s::%s' % (key, hash_args, hash_kw)             #5
```

The function `get_key` takes a function and its positional and keyword arguments (#1). We extract the module name and function name from the function (#2). Now we convert all positional arguments into a list of strings (#3). We convert the keyword arguments into a list of strings using the keys and the string representation of the values (#4). Finally, we return a string that consists of the three strings we have assembled so far (#5).

Now we use our function for a decorator to memoize (a term for the kind of caching we perform) previously calculated results:

```
# file: cache_deterministic.py
# form Ziade 2008
```

5 Caching

```
"""Example for a deterministic cache
"""

import functools

from get_key import get_key #1

cache = {} #2

def memoize_deterministic(get_key=get_key, cache=cache): #3
    """Parameterized decorator for memoizing.
    """

    def _memoize(function): #4
        """This takes the function.
        """

        @functools.wraps(function)
        def __memoize(*args, **kw): #5
            """This replaces the original function.
            """

            key = get_key(function, *args, **kw) #6
            try:
                return cache[key] #7
            except KeyError:
                value = function(*args, **kw) #8
                cache[key] = value #9
                return value #10

        return __memoize
    return _memoize
```

We use our function `get_key` (#1) and define a global dictionary that will be used to store pre-calculated data (#2). Our decorator takes the function and the dictionary as arguments (#3). This allows us to use other functions to retrieve a key and other caches possibly data dictionary-like data stores such as `shelve`. The second level function takes the function that is to be called as argument (#4). The third level function takes the arguments (#5). Now we retrieve our key (#6) and try to access the result from our cache (#7). If the key is not in the cache, we call our function (#8), store the result in the cache (#9) and return the result (#10).

Let's try how it works. We import the time modul and our module with the decorator:

```
>>> import time
>>> import cache_deterministic
```

We define a new function that adds to numbers and is decorated:

```
>>> @cache_deterministic.memoize_deterministic()
... def add(a, b):
```

5.3 Non-deterministic caching

```
...     time.sleep(2)
...     return a + b
...
```

We simulate some heavy calculations by delaying everything for two seconds with `sleep`. Let's call function:

```
>>> add(2, 2)
4
```

This took about two seconds. Do it again:

```
>>> add(2, 2)
4
```

Now the return is immediate.

Again:

```
>>> add(3, 3)
6
```

Two seconds delay. But now:

```
>>> add(3, 3)
```

Instantaneous response.

5.3 Non-deterministic caching

For non-deterministic caching, we use an age that the computed value should not exceed:

```
# file: cache_non_deterministic.py
# from Ziade 2008

"""Example for a cache that expires.
"""

import functools
import time

from get_key import get_key

cache = {}

def memoize_non_deterministic(get_key=get_key, storage=cache,
                              age=0): #1
    """Parameterized decorator that takes an expiration age.
```


5.3 Non-deterministic caching

```
"""

def _memoize(function):
    """This takes the function.
    """

    @functools.wraps(function)
    def __memoize(*args, **kw):
        """This replaces the original function.
        """
        key = get_key(function, *args, **kw)
        try:
            value_age, value = storage[key] #2
            deprecated = (age != 0 and #3
                          (value_age + age) < time.time())
        except KeyError: #4
            deprecated = True
        if not deprecated: #5
            return value
        storage[key] = time.time(), function(*args, **kw) #6
        return storage[key][1] #7
    return __memoize
return _memoize
```

This decorator is a variation of the deterministic one above. We can supply an age (#1). The value will be recalculated if this age is exceeded. We retrieve an age and a value from our cache (#2). The value will be deprecated, i.e. recalculated if we provide a non-zero age and the old age plus the specified age are smaller than the current time (#3). Note: This means, if you provide no age or an age of zero, the cache will never expire. The value will also be calculated if the key was not found (#4). We return the value if it is still valid (#5). Otherwise, we recalculate it and store it together with current time in the cache (#6) and return the freshly calculated value (#7).

Let's see how this works. We import our non-deterministic cache:

```
>>> import cache_non_deterministic
```

and define a new function with a maximum cache age of 5 seconds:

```
>>> @cache_non_deterministic.memoize_non_deterministic(age=5)
... def add2(a, b):
...     time.sleep(2)
...     return a + b
... 
```

The first call takes about two seconds:

```
>>> add2(2, 2)
4
```

5.4 Least Recently Used Cache

Immediately after this we do it again and get the response without delay:

```
>>> add2(2, 2)
4
```

Now we wait for at least 5 seconds ... and do it again:

```
>>> add2(2, 2)
4
```

This took again two seconds because the cache was not used and the value was recalculated due to the time since the last call being greater than five seconds.

5.4 Least Recently Used Cache

Caching is a common task. Therefore, Python starting from 3.2, provides a least recently used cache implementation in the `functools` module. This implementation is more elaborate than our attempts here. For example, it is thread-safe. There is a backport for Python 2.6 and higher. You can install it with:

```
pip install backports.functools_lru_cache
```

The size of the cache is determined by `maxsize`. There are three distinct case:

1. If `maxsize` is 0, there is no caching and the function result will be calculated every time the function is called.
2. If `maxsize` is `None`, it works similarly to our deterministic cache. That is, the LRU functionality is disabled and the cache can grow without limits.
3. If `maxsize` is a positive integer, the most recent `maxsize` function results are cached. This is the most interesting case because we actually use the LRU features.

Let's try it with a very small cache size of 2 to quickly see the effect of a filled cache:

```
import time
@lru_cache(maxsize=2)
def add(a, b):
    time.sleep(2) # wait two seconds
    return a + b
```

Now we use it:

```
>>> add(2, 2)
# takes two seconds
```

The second time around it returns without the two-second delay:

```
>>> add(2, 2)
# returns immediately
```

5.5 Memcached

Now we call the function with two other combinations of arguments:

```
>>> add(10, 5)
# takes two seconds
>>> add(10, 20)
# takes two seconds
```

Now our small cache of size 2 is filled with two other results and a call with our original combination will delay again for the first call:

```
>>> add2(2, 2)
# takes two seconds
>>> add2(2, 2)
# returns immediately
```

We can look at the original function:

```
>>> add.__wrapped__
<function __main__.add>
```

as well as at the cache:

```
>>> add.cache_info()
CacheInfo(hits=6, misses=3, maxsize=2, currsize=2)
```

Clearing the cache sets everything back to zero:

```
>>> add.cache_clear()
```

```
>>> add.cache_info()
CacheInfo(hits=0, misses=0, maxsize=2, currsize=0)
```

It is recommended to use `functools.lru_cache` over own solutions wherever possible. It is likely to be more stable. For example, coming up with your own solution that is thread-safe might be not as simple as it seems at first glance.

5.5 Memcached

[Memcached](#)⁵ is a caching server that is primarily used to speed up database based dynamic web pages. It is very fast, trades RAM for speed, and is very powerful. We don't have time to look at it here. There are several ways to use Memcached from Python. Also the probably most popular web framework Django uses Memcached ([Djangos cache](#)⁶).

6 Compilation of Tools for Speedup and Extending

There are many more ways to extend Python. Therefore, a short compilation of methods and tools for this purpose is given here. The compilation is by no means exhaustive.

Method/Tool	Remarks	Link
algorithmic improvements	try this first	http://www.python.org
NumPy	matlab like array processing	http://numpy.scipy.org
PyPy	fast Python implementation	http://pypy.org/
Cython	C with Python syntax	http://cython.org/
ctypes	call DLLs directly	Python's standard library
cffi	new interface to C for CPython and PyPy	https://cffi.readthedocs.org
Numba	compile to LLVM code	http://numba.pydata.org/
f2py	stable Fortran extensions	http://cens.ioc.ee/projects/f2py2e
C extensions by hand	lots of work	http://docs.python.org/ext/ext.html
SWIG	mature, stable, widely used	http://www.swig.org
Boost.Python	C ++ template based, elegant	http://www.boost.org/libs/python/doc
SIP	developed for Qt, fast	http://www.riverbankcomputing.co.uk/sip
PyInline	inline other languages (alpha)	http://pyinline.sourceforge.net
Theano	mathematical expressions (GPU)	http://deeplearning.net/software/theano/
PyCUDA	Python on GPGPUs	http://mathema.tician.de/software/pycuda
PyOpenCL	Python on GPGPUs	http://mathema.tician.de/software/pyopencl
Copperhead	Python on GPGPUs	http://code.google.com/p/copperhead/
COM/DCOM, CORBA, XML-RPC, ILU	middleware	various
Weave	inline C++ in Python	http://scipy.org/Weave
Babel	unite C/C++, F77/90, Py, Java	http://www.llnl.gov/CASC/components
pyufora	compiled, automatically parallel Python for data-science	http://docs.pyfora.com/en/stable/

7 End

7.1 Colophon

This material is written in `reStructuredText` and has been converted into PDF using `rst2pdf`. All Python files are dynamically included from the development directory.

7.2 Links

LANG2006	Hans Petter Lantangen, Python Scripting for Computational Science, Second Edition, Springer Berlin, Heidelberg, 2006.
MART2005	Alex Martelli et al., Python Cookbook, O'Reilly, 2nd Edition, 2005.
MART2006	Alex Martelli, Python in a Nutshell, O'Reilly, 2nd Edition, 2006.
ZIAD2008	Tarek Ziadè, Expert Python Programming: Best practices for designing, coding, and distributing your Python software, Packt 2008.
1	http://aspn.activestate.com/ASPN/Python/Cookbook
2	http://jiffyclub.github.io/snakeviz/
3	http://packages.python.org/Pympler/
4	http://en.wikipedia.org/wiki/Big_O_notation
5	http://www.danga.com/memcached/
6	http://docs.djangoproject.com/en/dev/topics/cache/